



# Evaluating LLMs to Measure Psychiatric Functioning

**Group Members:** Hannah Park, George Zhou, Jin Choi, Stephen Spivack  
**Mentors:** Matteo Malgaroli, Kyunghyun Cho, Geoffrey M. Reed

## RESEARCH DESCRIPTION

### Project Goals

- Assess the accuracy of Large Language Models (LLMs) to infer ICD-11 psychiatric diagnoses from clinically validated case vignettes.
- Assess the diagnostic accuracy of LLMs in evaluating psychiatric diagnosis across multiple languages
- Study agreement between LLMs responses and diagnostic assessments made by a diverse group of mental health professionals

### Research Approach

- Design LLMs pipeline to systematically feed data and collect response.
- Compare model performance of different LLMs.
- Improve diagnostic accuracy by experimenting different prompting strategies and through attention weight analysis.

### Related Work:

- Clinical Vignettes validation studies performed by World Health Organization (WHO). It was evaluated by 16,000 mental health and primary care professionals across 158 countries

## DATASET AND METRICS

### Data

- Data from three studies funded by the World Health Organization on Anxiety, Stress, and Mood disorders.
- Data collection through Qualtrics surveys encompassed four essential components:
  - Guidelines of Mental Disorder Diagnosis (ICD-11)
  - Clinical vignettes
  - Clinicians' responses to vignettes
  - Ground truth for validation
- Surveys were provided in multiple languages: English, Spanish, Chinese French, German, Russian, and Japanese.

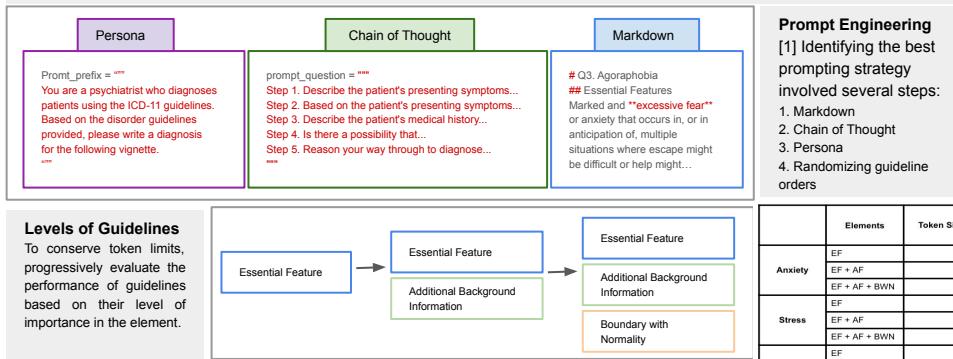
### Evaluation Metrics

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

## METHODS

### Model Capacity

The main challenge of the project was the large token size which exceeded the LLM input limit. Several methods were implemented to tackle this issue, which were to (1) use LLMs that take much larger inputs, such as Long-Llama, Llama2-32k, or Mistral, (2) summarize the ICD-11 guidelines, (3) incorporate 4-bit quantization, and (4) work with experts and collaborators to find which element are most important in the guideline.



## RESULTS

Generally, the Mistral model performs better than the Llama2 model except for Mood. Typically feeding in more information improves the model results for Mistral and Llama2 for Stress vignettes. However, for anxiety the models performs worse except when adding chain of thought and persona.

Interestingly, when the models get the wrong diagnosis, they tend to pick diagnosis that are **emblematic** of each disorder type. For anxiety it is GAD, stress is PTSD, and mood is single depressive disorder. When the correct diagnosis is no diagnosis, the model will usually give a false positive. There **were no false negative results**.

Prompt Type	Vignettes		
	Anxiety	Stress	Mood
Llama 2+EF	0.182	0.09	0.238
Llama 2+EF+AF	NA	0.182	0.2
Llama 2+EF+AF+BWN	NA	0.273	0.409
Mistral+EF	0.364	0.272	0.19
Mistral+EF+AF	0.273	0.273	0.19
Mistral+EF+A+BWN	0.182	0.455	0.19

Table 1: Model Performance with Zero-Shot Prompting

Prompt Type	Vignettes		
	Anxiety	Stress	Mood
Mistral+EF+A+BWN+CoT	0.455	0.273	0.19
Mistral+EF+A+BWN+CoT+Persona	0.545	0.545	0.333

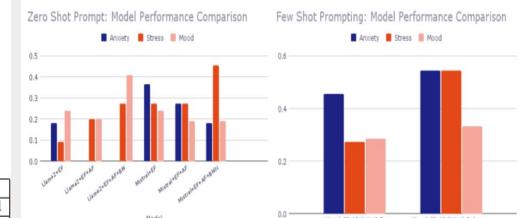
Table 2: Model Performance with Few-Shot Prompting

## CONCLUSION

The project aimed to identify the most effective ways to improve clinical-decision making rather than simply achieving correct diagnoses. Our models' behavior closely mirrored the typical clinician responses including misdiagnoses typically seen in clinical practice. Additional information didn't necessarily improve diagnostics, indicating that tailored prompting strategies for different disorders might optimize accuracy. This research is a step towards a more nuanced approach of leveraging AI for medical diagnosis.

## FUTURE PLANS

- Translate the data currently in different languages to English and format it according to our existing prompt formats for the purpose of linguistic comparisons.
- Examine the **inter-rater agreement between LLMs and clinicians survey responses** to explore similarities in mistakes and response patterns.
- Lastly, **study potential bias** that can affect diagnostic abilities of the model such as varied socio-demographic characteristics amongst patients.



Figures 1 and 2. Comparison of Prompting Strategies

**References:** [1] Perez, Ethan, et al. "True Few-Shot Learning with Language Models: 35th Conference on Neural Information Processing Systems, NeurIPS 2021, Advances in Neural Information Processing Systems 34 - 35th Conference on Neural Information Processing Systems, NeurIPS 2022, edited by Marc'Aurelio Ranzato et al., 2021, pp. 11054–70. NYU Scholars, <http://www.scopus.com/inward/record.uri?scp=85131820022&partnerID=8YFLogK>